ParGA: A Parallelized Genetic Algorithm Framework in MATLAB

William McGinnis

Abstract—**Genetic Algorithms (GAs) are frequently used to find solutions to optimization and search problems. They find solutions to these problems using techniques based on the concepts of natural evolution, namely mutation, inheritances and selection. The mechanics of the algorithm itself lends it to parallelization, which can speed it up by huge amounts, even on standard computers. ParGA leverages the MATLAB parallel computing toolbox to parallelize a real-coded, single crossover GA. It's modular design allows an objective function and fitness function to be simply plugged in.**

I. FRAMEWORK

ParGA is a MATLAB based framework for parallelized optimization problems. It uses a real-coded genetic algorithm, in which the function calls for determining population fitness are parallelized. The code is designed using the object oriented tools in MATLAB to allow for simple modification. Exact usage is detailed in the Case Study presented later in the report.

The details of the algorithm itself pertain mostly to it's style of breeding, mutation, evolution and the parallelization itself.

A. Breeding

The algorithm uses single crossover breeding. Future versions will include n-order crossover, where the user can select 1,2 or n crossovers. In each breeding, the most fit half of the generation continues on unmodified into the next generation. The 2nd half of the new generation is formed by using single crossover breeding on random members of the last population, with those which are most fit being most likely to breed.

After this is done, the new generation is mutated at a user supplied rate and by a user supplied amount. The mutation is performed with the following formula:

```
1  weights(index)=weights(index)*(((randn-1)/
       obj.kurt)+1);
```

Where obj.kurt is the user supplied kurtosis value. The larger it is, the smaller the mutations will be.

B. Neighbourhooding

Greater variation and the opportunity for multiple optima are accomplished by implementing neighbourhooding into the algorithm. This means that N discrete, separate populations are evolved for some amount before being combined into one final population. The amount of populations can be user selected, but care must be taken to ensure that the number of members per population is divisible by the number of neighbourhoods, otherwise there will be an error when merging the neighbourhoods.

C. Parallelization

While the framework will run without error in serial, if the user has the MATLAB Parallel Computing Toolbox, then simply adding "matlabpool" into the run script, or typing it into the command window before running an optimization will let the code run in parallel. This greatly speeds up the optimizations, especially when the fitness function is very computationally expensive.

The code can also be run on more powerful machines such as clusters, with the configuration requirements depending on the cluster itself.

Generally, the ideal number of cores to have is however many members per population you have, plus one for overhead. Any more will not help very much if at all.

II. CASE STUDY: SOLID ROCKET MOTOR DESIGN

In this case study, a solid rocket motor (SRM) legacy code was acquired from Dr. Hartfield, which was to be used to design a SRM to either have a specific thrust or chamber pressure output.

The desired output for thrust is a neutral 60,000 UNITS, and the desired output for chamber pressure is a neutral 500 UNITS.

A. FORTRAN Wrapper Function

The legacy code was written in FORTRAN and used input and output files, which posed a small problem. First and most importantly, input and output files cannot be used in the parallelized GA, because many nodes will be calling the function simultaneously, and the files cannot be read by multiple nodes simultaneously. To solve this, the program was modified to run from command line.

This is done in FORTRAN by using the IARGC() and GETARG() commands. IARGC() returns the number of parameters passed through command line, and GETARG() returns those parameters. Say, for instance, that you wanted to pass one variable, rpvar, to a program through command line. To do this, you can use a code similar to the one below to get the arguments.

```fortran
       integer :: jcount
       real :: jtemp
       character(len=32) :: arg

       jcount=IARGC()
       if (jcount==1) then
           call GETARG(1,arg)
           read(arg,*)jtemp
           rpvar=DBLE(jtemp)
       end if
```

The variable, rpvar, can then be used to compute some kind of fitness (to be minimized), and the fitness is printed to command line with the code:

```fortran
       print*, "", fitness
```

This will write the value stored in the variable, fitness, to the command line, where it can be read by another program.

B. MATLAB Integration

In the case of the SRM code, the FORTRAN wrapper function takes in 8 variables, uses them to calculate a thrust or chamber pressure profile, and returns the percent RMS error. In both cases, the RMS error is to be minimized, with a goal of under 1%. As with any optimization problem solved with ParGA, the first step is to modify the trigger function in Member.m. This function should contain all code to calculate the fitness of the Member, given the gene, stored in obj.gene.

The entire trigger function in this case is:

```matlab
function out=trigger(obj)
    command=sprintf('SRM_serve.exe %20.10
        fd0 %20.10fd0 %20.10fd0 %20.10fd0
        %20.10fd0 %20.10fd0 %20.10fd0
        %20.10fd0', obj.gene);
    [~,result]=system(command);
    out=str2double(result);
end
```

Next, the run script is written to actually build and perform the algorithm.

```matlab
close all; clear all;clc;

matlabpool;      %open local

mins=[0.01,0.01,.1,0.01,0.01,
    0.01,0.1,0.01];

maxes=[0.95,0.99,13,0.2,
    0.999999,0.77,100,0.667];

rate=0.05;
kurt=10;

gen_sample1=Globe(4,1000,8,mins,maxes,kurt
    ,rate);

gen_sample2=gen_sample1.evolveCommunities;

gen_sample3=gen_sample2.mergeCommunities;

gen_sample4=gen_sample3.evolveGlobe;
```

This will create an optimization with 4 neighborhoods, 1000 generations per epoch, 8 members per population, a mutation rate of 5%, and a kurtosis parameter of 10. Line 14 evolves the 4 seperate neighborhoods 1000 gerenations each. Line 16 takes the best quarter of each neighborhood and combines them into one population. Finally, line 18 evolves the combined population for 1000 generations. The final population can be evolved again as many times as desired by calling gen_sample4.evolveGlobe again.

Once this is completed, the final solution can be printed with the code:

```matlab
popfitness=gen_sample4.nations.popfitness;
[empty,ix]=sort(popfitness);
gene=gen_sample4.nations.pop{1,ix(1)}.gene
    ;
command=sprintf('SRM_serve_print.exe
    %20.10fd0 %20.10fd0 %20.10fd0 %20.10
    fd0 %20.10fd0 %20.10fd0 %20.10fd0
    %20.10fd0',gene);
[~,result]=system(command);
plot_data=sscanf(result,'%f\t');
```

```
 8 | figure;
 9 | plot(plot_data);
10 | ylabel('Thrust');
11 | title('Neutral Thrust Result')
```

Where the program, SRM_serve_print.exe, returns the entire thrust time series.

C. Results

The optimization was run on a quad core desktop computer with the parameters described in the previous section, for thrust. This amounts to a total of 40,000 function calls, which when parallelized is equivalent to 10,000 function calls in serial. The resultant RMS error percentage was 0.0001277538%, which was far better than the goal of 1%. The gene which resulted in this error was:

$$
\begin{aligned}
rpvar &= 0.5519 \\
rivar &= 0.3042 \\
xnx &= 2.1415 \\
fvar &= 0.0867 \\
epsx &= 0.0940 \\
fnlx &= 0.0314 \\
glx &= 21.3500 \\
diathx &= 0.6670
\end{aligned}
$$

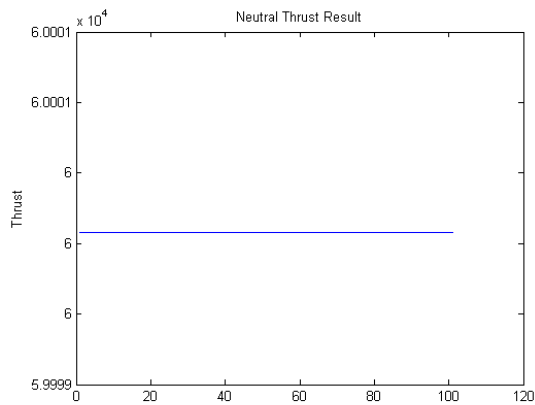And the resultant thrust profile is shown in Figure 1.

Figure 1.   Thrust Profile of Optimal Solution

This result was found in 819.71 seconds. This amounts to a 1.84x speedup over the time to run in serial of 1514.24 seconds.

III. CODE

All of the code for this toolbox can be found at:

https://github.com/wdm0006/ParGA